

STAT 243 Final Group Project: Adaptive Rejection Sampling

Jinze Gu, Havard Kvamme, Reid Stevens and Yang Wang

December 13, 2013

The R file was submitted by Jinze Gu.

Introduction

Adaptive Rejection Sampling (ARS) is a technique for generating random samples from a distribution with a log-concave density function developed by Gilks and Wild (1992). Given a set of starting values, ARS constructs piecewise-linear upper bounds using tangent lines to the log density function and piecewise-linear lower bounds using secant lines to the density function. The bounds are used to determine whether sampled points are accepted or rejected. These bounds are improved during the sampling process using the chosen values, which reduces the probability of needing to evaluate the density function.

Overview

Our ‘ars’ function performs sampling from a log-concave probability density function using ARS. The inputs to the ‘ars’ function are the log-concave density function of interest (‘f’), the number of points to sample from the density function (‘n’), and the initial set of abscissae (‘x_init’). The left and right bounds on the domain of the function (‘left_bound’ and ‘right_bound’) are assumed to be $\pm\infty$ unless specified. Extra arguments can also be passed to the function through the ‘ars’ function.

The ‘ars’ function is composed of seven subroutines: ‘make_z’, ‘make_upper_bound’, ‘make_lower_bound’, ‘sample_upper_bound’, ‘filter’, ‘update_sample’, and ‘update_x’. The initial step takes the set of starting values (‘x_init’) and evaluates the log of the density function at these points (creating the vector ‘hx’). The derivative of the log of the density function is estimated at the initial values with the auxiliary function ‘numericDeriv’ (creating the vector ‘hpx’). If the density function is unbounded, the function checks that the derivative at the smallest initial value is positive, while the derivative at the largest initial value is negative. If this condition is not met, an error is returned. The ‘hx’ and ‘hpx’ vectors are passed to the ‘make_z’ function.

Overall Tests: We also include a number of tests for the overall ‘ars’ function. The first test checks that the function will catch inputs that do not satisfy the requirements of the ARS algorithm (i.e., the derivatives of the first (last) initial points must be positive (negative)). The second test plots the sampled quantiles from a normal distribution against the theoretical values, to ensure the sampled values are correct. The third test confirms that extra arguments can be passed to the function. The fourth test confirms that the function can handle a bounded distribution. The fifth test checks that the function can catch a set of initial points that are too close together.

Make_z

Description: The ‘make_z’ function calculates the point of intersection between tangent lines to the log density function at successive abscissae x. The

‘make_z’ function takes as inputs the initial set of points, the left and right bounds, and the derivatives of the log density function at these points. The ‘make_z’ function uses equation (1) from Gilks and Wild (1992) to estimate the intersection points, using the lower and upper bounds as the first and last intersection points, respectively. These intersection points are used to compute the intervals over which each tangent line is the tightest upper bound of the log density function in the ‘make_upper_bound’ function. Though each tangent line is an upper bound over the domain of the entire log-density function, the ‘z’ vector allows us to identify the tangent line in each segment of the domain that is the tightest upper bound.

Tests: The ‘make_z’ function is tested with a plot of the ‘z’ values along with the associated values of a log-concave density function. Since the ‘z’ are just points from the density function, which is lognormal in this test case, we plot the two sets of values together to and see that they fall on the same line.

Make_upper_bound and Make_lower_bound

Description: The ‘make_upper_bound’ and ‘make_lower_bound’ functions generate the piecewise-linear upper and lower bounds to the log density function over the entire domain. The ‘make_upper_bound’ function simply constructs the piecewise linear upper bound using a different tangent line at the intervals given by the ‘z’ vector (which is the output of ‘make_z’). The output of the function (‘upper_bound’), give parameters for each linear section of the upper bound, which are simply the point of tangency (‘hx’) and the slope of the tangent line (‘hpx’) at each of the intervals given by the ‘z’ vector. The ‘make_lower_bound’ function returns a piecewise-linear lower bound for log-density (‘lower_bound’) using secant lines between successive abscissa (‘xx’) and their associated log density values (‘hx’). When an argument is out of the range of the abscissas of x , the lower bound is set to $-\infty$. The remaining functions deal with sampling from the distribution using the upper and lower bounds created by these two functions.

Tests: The ‘make_upper_bound’ and ‘make_lower_bound’ functions are tested by first confirming that all lower bounds are less than the corresponding upper bounds. We then test these functions by plotting the upper and lower bounds of a simple function ($f(x) = -x^2$). This test allows us to visually inspect whether the upper bounds are constructed correctly via tangent lines (for the upper bounds) and the secant lines (for the lower bounds).

Sample_upper_bound

Description: The ‘sample_upper_bound’ function takes as inputs ‘x’, ‘hx’, ‘hpx’, and ‘z’ and outputs a sample of ‘m’ points from the upper bound function. The points are sampled via the inverse CDF method, which is estimated analytically and implemented in the auxiliary function “inversecdf”. First, the

normalized integral of the upper bound function is estimated at each section of the piecewise function. Then ‘m’ points are sampled from a uniform distribution between zero and the ‘Inormalize’ value, where the ‘Inormalize’ value is the sum of the normalized integral of the upper bound taken at each interval. The interval each sample falls in is calculated (‘sample_interval’), and the inverse CDF is computed (using the ‘inversecdf’ function) at those points. The ‘inversecdf’ function returns a set of candidate points (‘cand’) that are passed to the ‘filter’ function.

Test: The ‘sample_upper_bound’ is tested by plotting a histogram of the values generated by the ‘sample_upper_bound’ function. These can be compared to the original function (in this case, $f(x) = -x^2$) to confirm that the sampled values match the function.

Filter and Update_sample

Description: The ‘filter’ takes in the candidate points (‘cand’) generated by the ‘sample_upper_bound’ function. The sample values are accepted if the associated randomly generated samples from the uniform distribution lie between the lower and upper bounds (i.e., the “squeeze test”). These points from the ‘cand’ vector are entered into the ‘accepted’ vector until the “squeeze test” is failed. From that point on, the log density function is evaluated and the value is accepted if the value lies between the upper bound and the density function (these points from ‘cand’ are entered into the ‘update’ vector). All other points are rejected. The ‘filter’ function returns a list that contains an index for the last accepted point and a value for the first non-accepted point. If all of the points in the ‘cand’ vector are accepted, the value for the first non-accepted point is simply “NA”. Likewise, if no points from the ‘cand’ vector are accepted, the first accepted point value is simply “NA”. The ‘update_sample’ function adds the accepted sample into the final sample and checks whether the updated sample should be added to the final sample as well. In addition, the log-concavity of these points is checked by confirming that log of the function values falls between the upper and lower bounds.

Tests: The ‘filter’ function is tested by inputting a sample from the upper bound (based on random samples from the normal distribution), to confirm that ‘filter’ properly separates these candidate values into the ‘accepted’, and the ‘non-accepted’ categories. The ‘update_sample_function’ is tested to ensure that, given candidate, accepted, and updated values, the function will return the index of the last accepted point. We also check that if no points are accepted, it returns 1 if the point satisfies the test “flag” it returns 0 if the point does not satisfy “flag.”

Update_x

Description: Finally, the ‘update_x’ function adds the points that are accepted

by the ‘filter’ function into the abscissae vector. The function also tests whether the updated sample is log-concave by checking whether the derivative at the updates points lie between their neighbouring points. An error is returned if the vector of derivatives at the updated points fail this test of log-concavity. The ‘while loop’ continues sampling, with the updated ‘x’ vector used in place of ‘x_init’ vector, until ‘n’ samples are accepted and the ‘ars’ function returns final accepted sample.

Tests: The ‘update_x’ test checks that the function output includes the correct points, with the associated function values and derivatives. Log-concavity is tested within the function by checking that the ‘hpx’ vector is non-decreasing. An error will be generated if the vector fails this test.

Contributions: Havard Kvamme worked on the filter, update_filter, make_lower/upper_bound functions. In addition, he worked on improving the efficiency of the functions, especially the make_lower/upper_bound. Jinze Gu wrote the first versions of make_z make_upper/lower bound, and h_prime functions using numericDeriv. Jinze also wrote the test function for each modular as well as the major ars function.

References:

Gilks, W. R. and Wild, P. (1992), “Adaptive Rejection Sampling for Gibbs Sampling,” *Applied Statistics*, Vol. 41, Issue 2, pp. 337-348.